


Introduction to MPI
2021 HPC Workshop: Parallel Programming

Alexander B. Pacheco

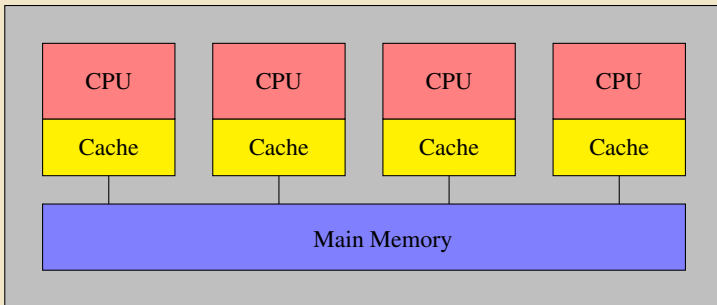
Research Computing

July 13 - 15, 2021


- Each process has its own address space
 - Data is local to each process
- Data sharing is achieved via explicit message passing
- Example
 - MPI



- All threads can access the global memory space.
- Data sharing achieved via writing to/reading from the same memory location
- Example
 - OpenMP
 - Pthreads



- The shared memory model is most commonly represented by Symmetric Multi-Processing (SMP) systems
 - Identical processors
 - Equal access time to memory
- Large shared memory systems are rare, clusters of SMP nodes are popular.



Shared Memory

- Pros
 - Global address space is user friendly
 - Data sharing is fast
- Cons
 - Lack of scalability
 - Data conflict issues

Distributed Memory

- Pros
 - Memory scalable with number of processors
 - Easier and cheaper to build
- Cons
 - Difficult load balancing
 - Data sharing is slow

- **Standardization:** MPI is the only message passing library that can be considered a standard. It is supported on virtually all HPC platforms.
- **Portability:** There is little or no need to modify your source code when you port your application to a different platform.
- **Performance Opportunities:** Vendor implementations should be able to exploit native hardware features to optimize performance. Any implementation is free to develop optimized algorithms.
- **Functionality:** There are over 430 routines defined in MPI-3, which includes the majority of those in MPI-2 and MPI-1.

Most MPI programs can be written using a dozen or less routines

- **Availability:** A variety of implementations are available, both vendor and public domain.

- MPI defines a standard API for message passing
 - The standard includes
 - What functions are available
 - The syntax of those functions
 - What the expected outcome is when calling those functions
 - The standard does NOT include
 - Implementation details (e.g. how the data transfer occurs)
 - Runtime details (e.g. how many processes the code run with etc.)
- MPI provides C/C++ and Fortran bindings

- There are two different types of MPI implementations commonly used.
 - ① **MPICH**: Developed by Argonne National Laboratory.
 - default MPI on Hawk.
 - used as a starting point for various commercial and open source MPI libraries
 - **MVAPICH2**: Developed by D. K. Panda with support for InfiniBand, iWARP, RoCE, and Intel Omni-Path.
default MPI on Sol.
 - **Intel MPI**: Intel's version of MPI.
available as part of the Intel OneAPI HPC Toolkit
 - **IBM MPI**: for IBM BlueGene, and
 - **CRAY MPI**: for Cray systems.
 - ② **OpenMPI**: A Free, Open Source implementation from merger of three well know MPI implementations.
 - **FT-MPI** from the University of Tennessee,
 - **LA-MPI**: from Los Alamos National Laboratory,
 - **LAM/MPI**: from Indiana University
 - Can be used for commodity network as well as high speed network.
 - available on Sol and Hawk but only libraries and a couple of packages are available.

- There is no MPI compiler available to compile programs nor is there is a compiler flag.
- Instead, you need to build the MPI libraries for a particular compiler.
- You can use MVAPICH2 and MPICH on Sol
- You should only use MPICH on Hawk.
- Each of these builds provide `mpicc`, `mpicxx` and `mpif90` for compiling C, C++ and Fortran codes respectively that are wrapper for the underlying compilers

```
[alp514.sol](793): module load mvapich2
[alp514.sol](794): mpicc --show
/share/Apps/intel/2020/compilers_and_libraries_2020.3.275/linux/bin/intel64/icc --lmpi -I/share/Apps/lusoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mvapich2/2.3.4-wguydha/include -L/share/Apps/lusoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mvapich2/2.3.4-wguydha/lib -Wl,-rpath -Wl,/share/Apps/lusoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mvapich2/2.3.4-wguydha/lib
[alp514.sol](795): module load mpich
```

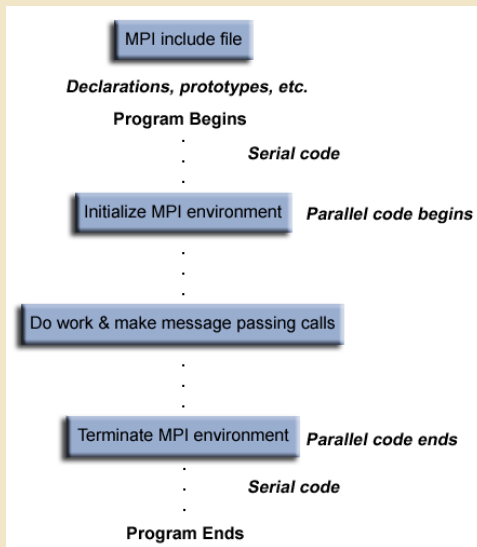
Lmod is automatically replacing "mvapich2/2.3.4" with "mpich/3.3.2".

```
[alp514.sol](796): mpif90 --show
/share/Apps/intel/2020/compilers_and_libraries_2020.3.275/linux/bin/intel64/iftor -L/share/Apps/lusoft/opt/spack/linux-centos8-haswell/intel-20.0.3/hwloc/2.2.0-rjrzfy7/lib -I/share/Apps/lusoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mpich/3.3.2-n7f36fo/include -I/share/Apps/lusoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mpich/3.3.2-n7f36fo/lib -lmpiifort -Wl,-rpath -Wl,/share/Apps/lusoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mpich/3.3.2-n7f36fo/lib -lmpi
```

- To run MPI applications, you need to launch the application using
 - `mpirun` (OpenMPI),
 - `mpirun_rsh` (MPICH and MVAPICH2), or
 - `mpiexec` (OpenMPI, MPICH and MVAPICH2).
- `mpirun`, `mpirun_rsh` and `mpiexec` are schedulers for the MPI library.
- On clusters with SLURM scheduler, `srun` can be used to launch MPI applications
- The MPI scheduler needs to be given additional information to correctly run MPI applications

	<code>mpiexec</code>	<code>mpirun_rsh</code>	<code>mpirun</code>
# Processors	<code>-n numprocs</code>	<code>-n numprocs</code>	<code>-np numprocs</code>
Processors List	<code>-hosts core1,core2,...</code>	<code>core1 core2 ...</code>	<code>-hosts core1,core2,...</code>
Processor filelist	<code>-f file</code>	<code>-hostfile file</code>	<code>-f/-hostfile file</code>

- Run an application myapp on 72 processors on a total of 3 nodes - node1, node2 and node3
 - `mpirun`: `mpirun -np 72 -f filename myapp`
 - `mpirun_rsh`: `mpirun_rsh -np 72 -hostfile filename myapp`
 - `mpiexec`: `mpiexec -n 72 -hosts node1,node2,node3 -ppn 24 myapp`
- The SLURM scheduler's `srun` launcher has information needed to run a mpi job
 - `srun`: `srun myapp`



- Header File: Required for all programs that make MPI library calls.


C	Fortran
<code>#include "mpi.h"</code>	<code>include 'mpif.h'</code> OR use <code>mpi</code>

- Format of MPI Calls:

- C names are case sensitive; Fortran names are not.
- Programs must not declare variables or functions with names beginning with the prefix MPI_ or PMPI_ (profiling interface)

C Binding	
Format	<code>rc = MPI_Xxxx(parameter, ...)</code>
Example	<code>rc = MPI_Bsend(&buf, count, type, dest, tag, comm)</code>
Error code	Returned as "rc". MPI_SUCCESS if successful
Fortran Binding	
Format	<code>call mpi_xxxx(parameter, ..., ierr)</code>
Example	<code>CALL MPI_BSEND(buf, count, type, dest, tag, comm, ierr)</code>
Error code	Returned as "ierr" parameter. MPI_SUCCESS if successful

- MPI uses objects called communicators and groups to define which collection of processes may communicate with each other.
- Most MPI routines require you to specify a communicator as an argument.
- **MPI_COMM_WORLD**: the default communicator contains all processes running a MPI program.



- Every process has its own unique, integer identifier assigned, called rank, by the system when the process initializes

- **MPI_INIT**: Initialize the MPI environment
- **MPI_COMM_SIZE**: Return total number of MPI processes
- **MPI_COMM_RANK**: Return rank of calling process
- **MPI_ABORT**: Terminates all MPI processes
- **MPI_GET_PROCESSOR_NAME**: Returns the processor name.
- **MPI_GET_VERSION**: Returns the version and subversion of the MPI standard
- **MPI_INITIALIZED**: Indicates whether MPI_Init has been called
- **MPI_WTIME**: Returns an elapsed wall clock time in seconds
- **MPI_WTICK**: Returns the resolution in seconds of MPI_WTIME
- **MPI_FINALIZE**: Terminate the MPI environment

C/C++

```
MPI_Init (&argc,&argv)
MPI_Comm_size (comm,&size)
MPI_Comm_rank (comm,&rank)
MPI_Abort (comm,errorcode)
MPI_Get_processor_name (&name,&resultlength)
MPI_Get_version (&version,&subversion)
MPI_Initialized (&flag)
MPI_Wtime ()
MPI_Wtick ()
MPI_Finalize ()
```

Fortran

```
MPI_INIT (ierr)
MPI_COMM_SIZE (comm,size,ierr)
MPI_COMM_RANK (comm,rank,ierr)
MPI_ABORT (comm,errorcode,ierr)
MPI_GET_PROCESSOR_NAME (name,resultlength,ierr)
MPI_GET_VERSION (version,subversion,ierr)
MPI_INITIALIZED (flag,ierr)
MPI_WTIME ()
MPI_WTICK ()
MPI_FINALIZE (ierr)
```


C

```
// required MPI include file
#include "mpi.h"
#include <stdio.h>

int main(int argc, char *argv[]) {
    int numtasks, rank, len, rc;
    char hostname[MPI_MAX_PROCESSOR_NAME];

    // initialize MPI
    MPI_Init(&argc,&argv);

    // get number of tasks
    MPI_Comm_size(MPI_COMM_WORLD,&numtasks);

    // get my rank
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);

    // print task number and rank
    printf ("Number of tasks= %d My rank= %d\n", numtasks,
        rank);

    // done with MPI
    MPI_Finalize();
}
```

Fortran

```
program simple

    implicit none
    ! required MPI include file
    include 'mpif.h'

    integer numtasks, rank, len, ierr
    character(MPI_MAX_PROCESSOR_NAME) hostname

    ! initialize MPI
    call MPI_INIT(ierr)

    ! get number of tasks
    call MPI_COMM_SIZE(MPI_COMM_WORLD, numtasks, ierr)

    ! get my rank
    call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)

    ! print task number and rank
    print '(a,i2,a,i2)', 'Number of tasks=',numtasks,' My
        rank=',rank

    ! done with MPI
    call MPI_FINALIZE(ierr)

end program simple
```

- Take the hello world code and add a few Environment Management functions
- Compile your code
- Run your code several different ways
- Examples to try out
 - 1 Print hostname
 - 2 Print mpi version
 - 3 Print hostname if your rank is odd and mpi version if rank is even

C

```
// required MPI include file
#include "mpi.h"
#include <stdio.h>

int main(int argc, char *argv[]) {
    int numtasks, rank, len, rc;
    char hostname[MPI_MAX_PROCESSOR_NAME];

    // initialize MPI
    MPI_Init(&argc,&argv);

    // get number of tasks
    MPI_Comm_size(MPI_COMM_WORLD,&numtasks);

    // get my rank
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);

    // get hostname
    MPI_Get_processor_name(hostname, &len);
    // print task number, rank and hostname
    printf ("Number of tasks= %d My rank= %d Running on %s\n",
           numtasks,rank,hostname);

    // done with MPI
    MPI_Finalize();
}
```

Fortran

```
program simple

    implicit none
    ! required MPI include file
    include 'mpif.h'

    integer numtasks, rank, len, ierr
    character(MPI_MAX_PROCESSOR_NAME) hostname

    ! initialize MPI
    call MPI_INIT(ierr)

    ! get number of tasks
    call MPI_COMM_SIZE(MPI_COMM_WORLD, numtasks, ierr)

    ! get my rank
    call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)

    ! get hostname
    call MPI_GET_PROCESSOR_NAME(hostname, len, ierr)

    ! print task number, rank and hostname
    print '(a,i2,a,i2,a,a)', 'Number of tasks=',numtasks,'
           My rank=',rank,&
           ' Running on ',hostname

    ! done with MPI
    call MPI_FINALIZE(ierr)

end program simple
```

```
[alp514.sol](1024): module load intel mpich
```

```
Lmod is automatically replacing ‘‘mvapich2/2.3.4’’ with ‘‘mpich/3.3.2’’.
```

```
[alp514.sol](1025): mpicc -o helloc hello.c
```

```
[alp514.sol](1026): mpif90 -o hellof hello.f90
```

```
[alp514.sol](1027): srun -p hawkgpu -n 4 -t 10 ./hellof
```

```
Number of tasks= 4 My rank= 0 Running on hawk-b624.cc.lehigh.edu
```

```
Number of tasks= 4 My rank= 1 Running on hawk-b624.cc.lehigh.edu
```

```
Number of tasks= 4 My rank= 3 Running on hawk-b624.cc.lehigh.edu
```

```
Number of tasks= 4 My rank= 2 Running on hawk-b624.cc.lehigh.edu
```

```
[alp514.sol](1028): srun -p hawkgpu -n 4 -t 10 ./helloc
```

```
Number of tasks= 4 My rank= 0 Running on hawk-b624.cc.lehigh.edu
```

```
Number of tasks= 4 My rank= 1 Running on hawk-b624.cc.lehigh.edu
```

```
Number of tasks= 4 My rank= 3 Running on hawk-b624.cc.lehigh.edu
```

```
Number of tasks= 4 My rank= 2 Running on hawk-b624.cc.lehigh.edu
```

```
[alp514.sol](1031): srun -p lts -n 4 -N 4 -t 10 ./helloc
```

```
Number of tasks= 4 My rank= 0 Running on sol-a105.cc.lehigh.edu
```

```
Number of tasks= 4 My rank= 2 Running on sol-a107.cc.lehigh.edu
```


```
Number of tasks= 4 My rank= 3 Running on sol-a108.cc.lehigh.edu
```

```
Number of tasks= 4 My rank= 1 Running on sol-a106.cc.lehigh.edu
```

- Not a part of the standard
 - Could vary from platform to platform
 - Or even from implementation to implementation on the same platform
 - mpicc/mpicxx/mpif77/mpif90: wrappers to compile MPI code and auto link to startup and message passing libraries



- Not a part of the standard
 - Could vary from platform to platform
 - Or even from implementation to implementation on the same platform
 - mpicc/mpicxx/mpif77/mpif90: wrappers to compile MPI code and auto link to startup and message passing libraries
- Unlike OpenMP and OpenACC, you cannot compile a MPI program for running in serial using the serial compiler
- The MPI program is not a standard C/C++/Fortran program and will spit out errors about missing libraries

- **MPI_COMM_WORLD**: the default communicator contains all processes running a MPI program.



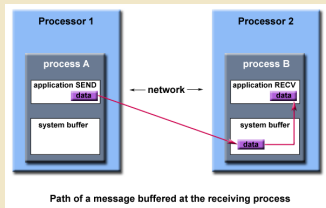
- Rank: unique id of each process
 - C: `MPI_Comm_Rank(MPI_Comm comm, int *rank)`
 - Fortran: `MPI_COMM_RANK(COMM, RANK, ERR)`
- Get the size/processes of a communicator
 - C: `MPI_Comm_Size(MPI_Comm comm, int *size)`
 - Fortran: `MPI_COMM_SIZE(COMM, SIZE, ERR)`

- Point-to-point communication functions
 - Message transfer from one process to another
- Collective communication functions
 - Message transfer involving all processes in a communicator



- MPI point-to-point operations typically involve message passing between two, and only two, different MPI tasks.
- One task is performing a send operation and the other task is performing a matching receive operation.
- There are different types of send and receive routines used for different purposes.
 - ① Blocking send / blocking receive
 - ② Non-blocking send / non-blocking receive
 - ③ Synchronous send

- Ideally, every send operation would be perfectly synchronized with its matching receive.
- MPI implementation must be able to deal with storing data when the two tasks are out of sync.
- Consider the following two cases:
 - ❶ A send operation occurs 5 seconds before the receive is ready - where is the message while the receive is pending?
 - ❷ Multiple sends arrive at the same receiving task which can only accept one send at a time - what happens to the messages that are backing up?
- MPI implementation (not the MPI standard) decides what happens to data in these types of cases.
- Typically, a system buffer area is reserved to hold data in transit.



System buffer space

- Opaque to the programmer and managed entirely by the MPI library
- A finite resource that can be easy to exhaust
- Often mysterious and not well documented
- Able to exist on the sending side, the receiving side, or both
- Something that may improve program performance because it allows send - receive operations to be asynchronous.

Blocking send / receive

- send will "return" after it is safe to modify the application buffer (your send data) for reuse
- send can be synchronous i.e. handshake with the receive task to confirm a safe send.
- send can be asynchronous if a system buffer is used to hold the data for eventual delivery to the receive.
- receive only "returns" after the data has arrived and is ready for use by the program.

Non-blocking send / receive

- behave similarly - they will return almost immediately.
- do not wait for any communication events to complete, such as message copying from user memory to system buffer space or the actual arrival of message.
- operations simply "request" the MPI library to perform the operation when it is able.

The user can not predict when that will happen.


- communications are primarily used to overlap computation with communication and exploit possible performance gains.

Order

- MPI guarantees that messages will not overtake each other.
- If a sender sends two messages (Message 1 and Message 2) in succession to the same destination, and both match the same receive, the receive operation will receive Message 1 before Message 2.
- If a receiver posts two receives (Receive 1 and Receive 2), in succession, and both are looking for the same message, Receive 1 will receive the message before Receive 2.
- Order rules do not apply if there are multiple threads participating in the communication operations.

Fairness

- MPI does not guarantee fairness - its up to the programmer to prevent operation starvation.



Blocking send / receive

- **MPI_Send:** Basic blocking send operation
- Routine returns only after the application buffer in the sending task is free for reuse.

```
MPI_Send (&buf, count, datatype, dest, tag, comm)
```

```
MPI_SEND (buf, count, datatype, dest, tag, comm, ierr)
```

- **MPI_Recv:** Receive a message
- will block until the requested data is available in the application buffer in the receiving task.

```
MPI_Recv (&buf, count, datatype, source, tag, comm, &status)
```

```
MPI_RECV (buf, count, datatype, source, tag, comm, status, ierr)
```


Non-blocking send / receive

- **MPI_Isend**: Identifies an area in memory to serve as a send buffer.
- Processing continues immediately without waiting for the message to be copied out from the application buffer

```
MPI_Isend (&buf, count, datatype, dest, tag, comm, &request)
```

```
MPI_ISEND (buf, count, datatype, dest, tag, comm, request, ierr)
```

- **MPI_Irecv**: Identifies an area in memory to serve as a receive buffer
- Processing continues immediately without actually waiting for the message to be received and copied into the application buffer

```
MPI_Irecv (&buf, count, datatype, source, tag, comm, &request)
```

```
MPI_IRECV (buf, count, datatype, source, tag, comm, request, ierr)
```

- **MPI_WAIT** and **MPI_TEST**: Functions required by nonblocking send and receive use to determine when the non-blocking receive operation completes and the requested message is available in the application buffer.

- **buf**: address space that references the data that is to be sent or received

In most cases, variable name that is to be sent/received

C programs: this argument is passed by reference and usually must be prepended with an ampersand: &var1

- **count**: number of data elements of a particular type to be sent
- **datatype**: MPI predefines its elementary data types

C type	MPI type
char	MPI_CHAR
unsigned char	MPI_UNSIGNED_CHAR
short	MPI_SHORT
unsigned short	MPI_UNSIGNED_SHORT
int	MPI_INT
unsigned int	MPI_UNSIGNED
long int	MPI_LONG
unsigned long int	MPI_UNSIGNED_LONG
long long int	MPI_LONG_LONG_INT
float	MPI_FLOAT
double	MPI_DOUBLE
long double	MPI_LONG_DOUBLE
unsigned char	MPI_BYTE

Fortran type	MPI type
character(1)	MPI_CHARACTER
integer	MPI_INTEGER
integer*2	MPI_INTEGER2
integer*4	MPI_INTEGER4
real	MPI_REAL
real*4	MPI_REAL4
real*8	MPI_REAL8
double precision	MPI_DOUBLE_PRECISION
complex	MPI_COMPLEX
double complex	MPI_DOUBLE_COMPLEX

- **dest**: indicates the process where a message should be delivered
specified as the rank of the receiving process
- **source**: indicates the originating process of the message.
specified as the rank of the sending process
- **tag**: arbitrary non-negative integer assigned by the programmer to uniquely identify a message
send and receive operations should match message tags
for a receive operation, the wild card **MPI_ANY_TAG** can be used to receive any message regardless of its tag.
- **comm**: indicates the communication context, or set of processes for which the source or destination fields are valid
unless the programmer is explicitly creating new communicators, the predefined communicator **MPI_COMM_WORLD** is usually used

- **status**: for a receive operation, indicates the source of the message and the tag of the message.

C: pointer to a predefined structure `MPI_Status`

Fortran: integer array of size `MPI_STATUS_SIZE`

- **request**: used by non-blocking send and receive operations

since non-blocking operations may return before the requested system buffer space is obtained, the system issues a unique request number

programmer uses this system assigned handle later determine completion of the non-blocking operation

C: a pointer to a predefined structure `MPI_Request`.

Fortran: an integer

- **MPI_Send**: Basic blocking send operation
- **MPI_Recv**: Receive a message and block until the requested data is available in the application buffer in the receiving task.
- **MPI_Ssend**: Synchronous blocking send: Send a message and block until the application buffer in the sending task is free for reuse and the destination process has started to receive the message
- **MPI_Sendrecv**: Send a message and post a receive before blocking.

Will block until the sending application buffer is free for reuse and until the receiving application buffer contains the received message

- **MPI_Probe**: Performs a blocking test for a message.

The wildcards `MPI_ANY_SOURCE` and `MPI_ANY_TAG` may be used to test for a message from any source or with any tag.

C: the actual source and tag will be returned in the status structure as `status.MPI_SOURCE` and `status.MPI_TAG`.

Fortran: they will be returned in the integer array `status(MPI_SOURCE)` and `status(MPI_TAG)`.

- **MPI_Get_Count**: Returns the source, tag and number of elements of datatype received.

Can be used with both blocking and non-blocking receive operations.

C: the actual source and tag will be returned in the status structure as `status.MPI_SOURCE` and `status.MPI_TAG`.

Fortran: they will be returned in the integer array `status(MPI_SOURCE)` and `status(MPI_TAG)`.

```
MPI_Ssend (&buf, count, datatype, dest, tag, comm)
MPI_Sendrecv (&sendbuf, sendcount, sendtype, dest, sendtag,
              &recvbuf, recvcount, recvtype, source, recvtag,
              comm, &status)
MPI_Probe (source, tag, comm, &status)
MPI_Get_count (&status, datatype, &count)
```

```
MPI_SSEND (buf, count, datatype, dest, tag, comm, ierr)
MPI_SENDRECV (sendbuf, sendcount, sendtype, dest, sendtag, &
              recvbuf, recvcount, recvtype, source, recvtag, &
              comm, status, ierr)
MPI_PROBE (source, tag, comm, status, ierr)
MPI_GET_COUNT (status, datatype, count, ierr)
```

- Modify the `pingpong.c` or `pingpong.f90` example to do a blocking send and receive.
- Task 0 sends a ping to task 1 and awaits return ping
- This example only requires two MPI processes
- What happens if you run on more than 2 cpus


```
[alp514.sol](1132): mpicc -o pingpong pingpong.c
[alp514.sol](1133): mpif90 -o pingpong pingpong.f90
[alp514.sol](1134): srun -n 2 -p hawkgpu -t 5 ./pingpong
Task 1: Received 1 char(s) from task 0 with tag 1
Task 0: Received 1 char(s) from task 1 with tag 1

[alp514.sol](1135): srun -n 2 -p hawkgpu -t 5 ./pingpong
Task 1 : Received 1 char(s) from task 0 with tag 1
Task 0 : Received 1 char(s) from task 1 with tag 1

[alp514.sol](1136): srun -n 4 -p hawkgpu -t 5 ./pingpong
Task 1 : Received 1 char(s) from task 0 with tag 1
Task 3 : Received 0 char(s) from task 0 with tag 0
Task 2 : Received 0 char(s) from task 0 with tag 0
Task 0 : Received 1 char(s) from task 1 with tag 1

[alp514.sol](1137): srun -n 4 -p hawkgpu -t 5 ./pingpong
Task 0: Received 1 char(s) from task 1 with tag 1
Task 1: Received 1 char(s) from task 0 with tag 1
Task 2: Received -32766 char(s) from task 2496 with tag -1075053569
Task 3: Received -32766 char(s) from task 1668810496 with tag 32588
```

- **MPI_Isend:** Identifies an area in memory to serve as a send buffer.
Processing continues immediately without waiting for the message to be copied out from the application buffer.
A communication request handle is returned for handling the pending message status.
The program should not modify the application buffer until the non-blocking send has completed.
- **MPI_Irecv:** Identifies an area in memory to serve as a receive buffer.
Processing continues immediately without actually waiting for the message to be received and copied into the the application buffer.
A communication request handle is returned for handling the pending message status.
The program must use calls to MPI_Wait or MPI_Test to determine when the non-blocking receive operation completes and the requested message is available in the application buffer.
- **MPI_Issend:** Non-blocking synchronous send.
Similar to MPI_Isend(), except MPI_Wait() or MPI_Test() indicates when the destination process has received the message.

- **MPI_Test[any,all,some]:** MPI_Test checks the status of a specified non-blocking send or receive operation.

The flag parameter is returned logical true (1) if the operation has completed, and logical false (0) if not.

For multiple non-blocking operations, the programmer can specify any, all or some completions.

- **MPI_Iprobe:** Performs a non-blocking test for a message.

The wildcards MPI_ANY_SOURCE and MPI_ANY_TAG may be used to test for a message from any source or with any tag. The integer flag parameter is returned logical true (1) if a message has arrived, and logical false (0) if not.

C: the actual source and tag will be returned in the status structure as status.MPI_SOURCE and status.MPI_TAG.

Fortran: they will be returned in the integer array status(MPI_SOURCE) and status(MPI_TAG).

- **MPI_Wait[any,all,some]:** MPI_Wait blocks until a specified non-blocking send or receive operation has completed.

For multiple non-blocking operations, the programmer can specify any, all or some completions.

```
MPI_Issend (&buf, count, datatype, dest, tag, comm, &request)
MPI_Test (&request, &flag, &status)
MPI_Testany (count, &array_of_requests, &index, &flag, &status)
MPI_Testall (count, &array_of_requests, &flag, &array_of_statuses)
MPI_Testsome (incount, &array_of_requests, &outcount,
              &array_of_offsets, &array_of_statuses)
MPI_Wait (&request, &status)
MPI_Waitany (count, &array_of_requests, &index, &status)
MPI_Waitall (count, &array_of_requests, &array_of_statuses)
MPI_Waitsome (incount, &array_of_requests, &outcount,
              &array_of_offsets, &array_of_statuses)
MPI_Iprobe (source, tag, comm, &flag, &status)

MPI_ISSEND (buf, count, datatype, dest, tag, comm, request, ierr)
MPI_TEST (request, flag, status, ierr)
MPI_TESTANY (count, array_of_requests, index, flag, status, ierr)
MPI_TESTALL (count, array_of_requests, flag, array_of_statuses, ierr)
MPI_TESTSOME (incount, array_of_requests, outcount, &
              array_of_offsets, array_of_statuses, ierr)
MPI_WAIT (request, status, ierr)
MPI_WAITANY (count, array_of_requests, index, status, ierr)
MPI_WAITALL (count, array_of_requests, array_of_statuses, ierr)
MPI_WAITSOME (incount, array_of_requests, outcount, &
              array_of_offsets, array_of_statuses, ierr)
MPI_IPROBE (source, tag, comm, flag, status, ierr)
```

- Modify the `ring.c` or `ring.f90` example to do a non blocking send and receive.
- Each process sends 1 to the left and 2 to the right


```
[alp514.sol](1160): mpicc -o ringc ring.c
[alp514.sol](1161): mpif90 -o ringf ring.f90
[alp514.sol](1162): srun -n 4 -p hawkgpu -t 10 ./ringc
Task 0: Received from task 3 with tag 1 and from task 1 with tag 2
Task 0: Send to task 3 with tag 2 and to task 1 with tag 1
Task 3: Received from task 2 with tag 1 and from task 0 with tag 2
Task 3: Send to task 2 with tag 2 and to task 0 with tag 1
Task 2: Received from task 1 with tag 1 and from task 3 with tag 2
Task 2: Send to task 1 with tag 2 and to task 3 with tag 1
Task 1: Received from task 0 with tag 1 and from task 2 with tag 2
Task 1: Send to task 0 with tag 2 and to task 2 with tag 1
```

```
[alp514.sol](1163): srun -n 4 -p hawkgpu -t 10 ./ringf
Task 0: Received from task 3 with tag 1 and from task 1 with tag 2
Task 0: Send to task 3 with tag 2 and to task 1 with tag 1
Task 2: Received from task 1 with tag 1 and from task 3 with tag 2
Task 2: Send to task 1 with tag 2 and to task 3 with tag 1
Task 3: Received from task 2 with tag 1 and from task 0 with tag 2
Task 3: Send to task 2 with tag 2 and to task 0 with tag 1
Task 1: Received from task 0 with tag 1 and from task 2 with tag 2
Task 1: Send to task 0 with tag 2 and to task 2 with tag 1
```

- Collective communication routines must involve all processes within the scope of a communicator.
- All processes are by default, members in the communicator `MPI_COMM_WORLD`.
- Unexpected behavior, including program failure, can occur if even one task in the communicator doesn't participate.
- It is the programmer's responsibility to ensure that all processes within a communicator participate in any collective operations.
- Collective communication routines do not take message tag arguments.

Types of Collective Operations

- **Synchronization:** processes wait until all members of the group have reached the synchronization point.
- **Data Movement:** broadcast, scatter/gather, all to all.
- **Collective Computation (reductions):** one member of the group collects data from the other members and performs an operation (min, max, add, multiply, etc.) on that data.



- **MPI_Barrier**: Creates a barrier synchronization in a group
- **MPI_Bcast**: Broadcasts (sends) a message from the process with rank root to all other processes in the group
- **MPI_Scatter**: Distributes distinct messages from a single source task to each task in the group
- **MPI_Gather**: Gathers distinct messages from each task in the group to a single destination task
- **MPI_Allgather**: Concatenation of data to all tasks in a group. Each task in the group, in effect, performs a one-to-all broadcasting operation within the group
- **MPI_Reduce**: Applies a reduction operation on all tasks in the group and places the result in one task
- **MPI_Allreduce**: equivalent to an **MPI_Reduce** followed by an **MPI_Bcast**
- **MPI_Reduce_scatter** equivalent to an **MPI_Reduce** followed by an **MPI_Scatter** operation
- **MPI_Alltoall**: Each task in a group performs a scatter operation, sending a distinct message to all the tasks in the group in order by index
- **MPI_Scan**: Performs a scan operation with respect to a reduction operation across a task group

```
MPI_Bcast (&buffer, count, datatype, root, comm)
MPI_Scatter (&sendbuf, sendcnt, sendtype, &recvbuf, recvcnt, recvtype, root, comm)
MPI_Gather (&sendbuf, sendcnt, sendtype, &recvbuf, recvcnt, recvtype, root, comm)
MPI_Allgather (&sendbuf, sendcount, sendtype, &recvbuf, recvcnt, recvtype, comm)
MPI_Reduce (&sendbuf, &recvbuf, count, datatype, op, root, comm)
MPI_Allreduce (&sendbuf, &recvbuf, count, datatype, op, comm)
MPI_Reduce_scatter (&sendbuf, &recvbuf, recvcnt, datatype, op, comm)
MPI_Alltoall (&sendbuf, sendcount, sendtype, &recvbuf, recvcnt, recvtype, comm)
MPI_Scan (&sendbuf, &recvbuf, count, datatype, op, comm)
```

```
MPI_BCAST (buffer, count, datatype, root, comm, ierr)
MPI_SCATTER (sendbuf, sendcnt, sendtype, recvbuf, recvcnt, recvtype, root, comm, ierr) MPI_GATHER (sendbuf, sendcnt,
    sendtype, recvbuf, recvcnt, recvtype, root, comm, ierr)
MPI_ALLGATHER (sendbuf, sendcount, sendtype, recvbuf, recvcnt, recvtype, comm, info)
MPI_REDUCE (sendbuf, recvbuf, count, datatype, op, root, comm, ierr)
MPI_ALLREDUCE (sendbuf, recvbuf, count, datatype, op, comm, ierr)
MPI_REDUCE_SCATTER (sendbuf, recvbuf, recvcnt, datatype, op, comm, ierr)
MPI_ALLTOALL (sendbuf, sendcount, sendtype, recvbuf, recvcnt, recvtype, comm, ierr)
MPI_SCAN (sendbuf, recvbuf, count, datatype, op, comm, ierr)
```

- Parallelize the `pi_mpi.f90` or `pi_mpi.c` making use of collective communication functions?

```
[alp514.sol](1061): mpicc -o pic pi_mpi.c
[alp514.sol](1062): mpif90 -o pif pi_mpi.f90
[alp514.sol](1063): srun -n 6 -p hawkgpu --reservation=lts_165 -t 10 ./pic
pi = 3.141592653589771
time to compute = 0.02951 seconds

[alp514.sol](1064): srun -n 6 -p hawkgpu --reservation=lts_165 -t 10 ./pif
pi = 3.141592633589724
time to compute = 0.023 seconds

[alp514.sol](1065): srun -N 2 --ntasks-per-node=3 -p infolab -t 10 ./pic
pi = 3.141592653589771
time to compute = 0.0282662 seconds

[alp514.sol](1066): srun -N 2 --ntasks-per-node=3 -p infolab -t 10 ./pif
pi = 3.141592633589724
time to compute = 0.015 seconds
```

```
[alp514.sol](1081): srun -N 2 --ntasks-per-node=12 -p chem -t 10 ./pic  
pi = 3.141592653589797  
time to compute = 0.014163 seconds
```

```
[alp514.sol](1082): srun -N 2 --ntasks-per-node=12 -p chem -t 10 ./pif  
pi = 3.141592633589827  
time to compute = 0.007 seconds
```

```
[alp514.sol](1083): srun -N 2 --ntasks-per-node=24 -p chem -t 10 ./pic  
pi = 3.141592653589789  
time to compute = 0.0264809 seconds
```

```
[alp514.sol](1084): srun -N 2 --ntasks-per-node=24 -p chem -t 10 ./pif  
pi = 3.141592633589773  
time to compute = 0.038 seconds
```

```
[alp514.sol](1085): srun -N 2 --ntasks-per-node=36 -p chem -t 10 ./pic  
pi = 3.141592653589792  
time to compute = 0.184978 seconds
```

```
[alp514.sol](1086): srun -N 2 --ntasks-per-node=36 -p chem -t 10 ./pif  
pi = 3.141592633589787  
time to compute = 0.014 seconds
```

• Books

- ➊ Parallel Programming with MPI by Peter Pacheco (No relation)
- ➋ Using MPI - 2nd Edition: Portable Parallel Programming with the Message Passing Interface (Scientific and Engineering Computation) by William Gropp
- ➌ Parallel Programming in C with MPI and Openmp by Michael J. Quinn
- ➍ MPI: The Complete Reference by Marc Snir *et. al.*
- ➎ Beginning MPI (An Introduction in C) by Wesley Kendall
Online version: <http://mpitutorial.com/>

• Tutorials

- ➊ MPI: <https://computing.llnl.gov/tutorials/mpi/>
- ➋ Advanced MPI:
<https://hpc.llnl.gov/sites/default/files/DavidCronkSlides.pdf>
- ➌ CITutor: <https://www.citutor.org/>
- ➍ XSEDE HPC Monthly Workshop Series:
<https://psc.edu/xsede-hpc-series-all-workshops>