

Introduction to MPI

2021 HPC Workshop: Parallel Programming


Alexander B. Pacheco

Research Computing


July 13 - 15, 2021

1 Introduction


- Each process has its own address space
 - Data is local to each process
- Data sharing is achieved via explicit message passing
- Example
 - MPI



- All threads can access the global memory space.
- Data sharing achieved via writing to/reading from the same memory location
- Example
 - OpenMP
 - Pthreads



- The shared memory model is most commonly represented by Symmetric Multi-Processing (SMP) systems
 - Identical processors
 - Equal access time to memory
- Large shared memory systems are rare, clusters of SMP nodes are popular.



Shared Memory

- Pros
 - Global address space is user friendly
 - Data sharing is fast
- Cons
 - Lack of scalability
 - Data conflict issues

Distributed Memory

- Pros
 - Memory scalable with number of processors
 - Easier and cheaper to build
- Cons
 - Difficult load balancing
 - Data sharing is slow

- There are already network communication libraries
- Optimized for performance
- Take advantage of faster network transport
 - Shared memory (within a node)
 - Faster cluster interconnects (e.g. InfiniBand)
 - TCP/IP if all else fails
- Enforces certain guarantees
 - Reliable messages
 - In-order message arrival
- Designed for multi-node technical computing

- MPI defines a standard API for message passing
 - The standard includes
 - What functions are available
 - The syntax of those functions
 - What the expected outcome is when calling those functions
 - The standard does NOT include
 - Implementation details (e.g. how the data transfer occurs)
 - Runtime details (e.g. how many processes the code run with etc.)
- MPI provides C/C++ and Fortran bindings

- There are two different types of MPI implementations commonly used.
 - ① **MPICH:** Developed by Argonne National Laboratory.
 - used as a starting point for various commercial and open source MPI libraries
 - **MVAPICH2:** Developed by D. K. Panda with support for InfiniBand, iWARP, RoCE, and Intel Omni-Path. (default MPI on Sol),
 - **Intel MPI:** Intel's version of MPI. Part of Intel OneAPI HPC Toolkit
 - **IBM MPI:** for IBM BlueGene, and
 - **CRAY MPI:** for Cray systems.
 - ② **OpenMPI:** A Free, Open Source implementation from merger of three well know MPI implementations. Can be used for commodity network as well as high speed network.
 - **FT-MPI** from the University of Tennessee,
 - **LA-MPI:** from Los Alamos National Laboratory,
 - **LAM/MPI:** from Indiana University

- ```
[alp514.sol](793): module load mvapich2
[alp514.sol](794): mpicc -show
/share/Apps/intel/2020/compilers_and_libraries_2020.3.275/linux/bin/intel64/icc -lmpi -I/share/Apps/luoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mvapich2/2.3.4-gwyudha/include -L/share/Apps/luoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mvapich2/2.3.4-gwyudha/lib -Wl,-rpath -Wl,/share/Apps/luoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mvapich2/2.3.4-gwyudha/lib
[alp514.sol](795): module load mpich
```

```
[alp514.sol](796): mpi-f90 --show
/share/Apps/intel/2020/compilers_and_libraries_2020.3.275/linux/bin/intel64/ifort -L/share/Apps/lufoft/opt/spack/linux-centos8-
haswell/intel-20.0.3/hwloc/2.2.0-rjrzfy7/lib -I/share/Apps/lufoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mpi/h/3.
3.2-n7f36fo/include -I/share/Apps/lufoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mpi/h/3.3.2-n7f36fo/include -L/
share/Apps/lufoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mpi/h/3.3.2-n7f36fo/lib -lmpifort -Wl,-rpath -Wl,/share/
App/lufoft/opt/spack/linux-centos8-haswell/intel-20.0.3/mpi/h/3.3.2-n7f36fo/lib -lmpi
```

- To run MPI applications, you need to launch the application using mpirun (OpenMPI), mpirun\_rsh (MPICH and MVAPICH2) or mpiexec (OpenMPI, MPICH and MVAPICH2).
- mpirun, mpirun\_rsh and mpiexec are schedulers for the MPI library.
- On clusters with SLURM scheduler, srun can be used to launched MPI applications
- The MPI scheduler needs to be given additional information to correctly run MPI applications

|                    | mpiexec                | mpirun_rsh      | mpirun                 |
|--------------------|------------------------|-----------------|------------------------|
| # Processors       | -n numprocs            | -n numprocs     | -np numprocs           |
| Processors List    | -hosts core1,core2,... | core1 core2 ... | -hosts core1,core2,... |
| Processor filelist | -f file                | -hostfile file  | -f/-hostfile file      |

- Run an application myapp on 72 processors on a total of 3 nodes - node1, node2 and node3
  - mpirun: `mpirun -np 72 -f filename myapp`
  - mpirun\_rsh: `mpirun_rsh -np 72 -hostfile filename myapp`
  - mpiexec: `mpiexec -n 72 -hosts node1,node2,node3 -ppn 24 myapp`

- ❶ Initiate communication between processes
  - **MPI\_INIT**: initialize MPI environment
  - **MPI\_COMM\_SIZE**: return total number of MPI processes
  - **MPI\_COMM\_RANK**: return rank of calling process
- ❷ Communicate data between processes
  - **MPI\_SEND**: send a message
  - **MPI\_RECV**: receive a message
- ❸ Terminate the MPI environment using **MPI\_FINALIZE**

## C

```
// required MPI include file
#include "mpi.h"
#include <stdio.h>

int main(int argc, char *argv[]) {
 int numtasks, rank, len, rc;
 char hostname[MPI_MAX_PROCESSOR_NAME];

 // initialize MPI
 MPI_Init(&argc,&argv);

 // get number of tasks
 MPI_Comm_size(MPI_COMM_WORLD,&numtasks);

 // get my rank
 MPI_Comm_rank(MPI_COMM_WORLD,&rank);

 // this one is obvious
 MPI_Get_processor_name(hostname, &len);
 printf ("Number of tasks= %d My rank= %d Running on %s\\n", numtasks,rank,hostname);

 // done with MPI
 MPI_Finalize();
}
```

## Fortran

```
program simple

! required MPI include file
include 'mpif.h'

integer numtasks, rank, len, ierr
character(MPI_MAX_PROCESSOR_NAME) hostname

! initialize MPI
call MPI_INIT(ierr)

! get number of tasks
call MPI_COMM_SIZE(MPI_COMM_WORLD, numtasks, ierr)

! get my rank
call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)

! this one is obvious
call MPI_GET_PROCESSOR_NAME(hostname, len, ierr)
print '(a,i2,a,i2,a,a)', 'Number of tasks=',numtasks,'
 My rank=',rank,&
 ' Running on ',hostname

! done with MPI
call MPI_FINALIZE(ierr)

end program simple
```

```
[alp514.sol](1003): module load mvapich2/2.2/intel-17.0.3
[alp514.sol](1004): mpicc -o hello hello.c
[alp514.sol](1005): mpif90 -o hellof hello.f90
[alp514.sol](1006): srun -p eng -n 4 ./hello
Number of tasks= 4 My rank= 3 Running on sol-b110
Number of tasks= 4 My rank= 2 Running on sol-b110
Number of tasks= 4 My rank= 1 Running on sol-b110
Number of tasks= 4 My rank= 0 Running on sol-b110
[alp514.sol](1007): srun -p eng -n 4 ./hellof
Number of tasks= 4 My rank= 3 Running on sol-b110
Number of tasks= 4 My rank= 2 Running on sol-b110
Number of tasks= 4 My rank= 0 Running on sol-b110
Number of tasks= 4 My rank= 1 Running on sol-b110
```

- Header File: Required for all programs that make MPI library calls.


| C                             | Fortran                       |
|-------------------------------|-------------------------------|
| <code>#include "mpi.h"</code> | <code>include 'mpif.h'</code> |

- Format of MPI Calls:

- C names are case sensitive; Fortran names are not.
- Programs must not declare variables or functions with names beginning with the prefix MPI\_ or PMPI\_ (profiling interface)

| C Binding       |                                                                                                      |
|-----------------|------------------------------------------------------------------------------------------------------|
| Format          | <code>rc = MPI_Xxxx(parameter, ... )</code>                                                          |
| Example         | <code>rc = MPI_Bsend(&amp;buf, count, type, dest, tag, comm)</code>                                  |
| Error code      | Returned as "rc". MPI_SUCCESS if successful                                                          |
| Fortran Binding |                                                                                                      |
| Format          | <code>CALL MPI_XXXXX(parameter,..., ierr)</code><br><code>call mpi_XXXXX(parameter,..., ierr)</code> |
| Example         | <code>CALL MPI_BSEND(buf, count, type, dest, tag, comm, ierr)</code>                                 |
| Error code      | Returned as "ierr" parameter. MPI_SUCCESS if successful                                              |

- A communicator is an identifier associated with a group of processes



```
MPI_Comm_size(MPI_COMM_WORLD, int &numtasks);
MPI_Comm_rank(MPI_COMM_WORLD, int &rank);
```

```
call MPI_COMM_SIZE(MPI_COMM_WORLD, numtasks, ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)
```



- A communicator is an identifier associated with a group of processes
  - Can be regarded as the name given to an ordered list of processes
  - Each process has a unique rank, which starts from 0 (usually referred to as "root")
  - It is the context of MPI communications and operations
    - For instance, when a function is called to send data to all processes, MPI needs to understand what "all"
  - `MPI_COMM_WORLD`: the default communicator contains all processes running a MPI program
  - There can be many communicators
    - e.g., `MPI_Comm_split(MPI_Commcomm, intcolor, int, kye, MPI_Comm* newcomm)`
  - A process can belong to multiple communicators
    - The rank is usually different

- Rank: unique id of each process
  - C: `MPI_Comm_Rank(MPI_Comm comm, int *rank)`
  - Fortran: `MPI_COMM_RANK(COMM, RANK, ERR)`
- Get the size/processes of a communicator
  - C: `MPI_Comm_Size(MPI_Comm comm, int *size)`
  - Fortran: `MPI_COMM_SIZE(COMM, SIZE, ERR)`

- Not a part of the standard
  - Could vary from platform to platform
  - Or even from implementation to implementation on the same platform
  - mpicc/mpicxx/mpif77/mpif90: wrappers to compile MPI code and auto link to startup and message passing libraries

- Not a part of the standard
  - Could vary from platform to platform
  - Or even from implementation to implementation on the same platform
  - mpicc/mpicxx/mpif77/mpif90: wrappers to compile MPI code and auto link to startup and message passing libraries
- Unlike OpenMP and OpenACC, you cannot compile a MPI program for running in serial using the serial compiler
- The MPI program is not a standard C/C++/Fortran program and will spit out errors about missing libraries



## ● Environment management functions

- 1 MPI\_INIT
- 2 MPI\_COMM\_SIZE
- 3 MPI\_COMM\_RANK
- 4 MPI\_ABORT: Terminates all MPI processes
- 5 MPI\_GET\_PROCESSOR\_NAME: Returns the processor name.
- 6 MPI\_GET\_VERSION: Returns the version and subversion of the MPI standard
- 7 MPI\_INITIALIZED: Indicates whether MPI\_Init has been called
- 8 MPI\_WTIME: Returns an elapsed wall clock time in seconds
- 9 MPI\_WTICK: Returns the resolution in seconds of MPI\_WTIME
- 10 MPI\_FINALIZE

```
MPI_Init (&argc,&argv)
MPI_Comm_size (comm,&size)
MPI_Comm_rank (comm,&rank)
MPI_Abort (comm,errorcode)
MPI_Get_processor_name (&name,&resultlength)
MPI_Get_version (&version,&subversion)
MPI_Initialized (&flag)
MPI_Wtime ()
MPI_Wtick ()
MPI_Finalize ()
```

```
MPI_INIT (ierr)
MPI_COMM_SIZE (comm,size,ierr)
MPI_COMM_RANK (comm,rank,ierr)
MPI_ABORT (comm,errorcode,ierr)
MPI_GET_PROCESSOR_NAME (name,resultlength,ierr)
MPI_GET_VERSION (version,subversion,ierr)
MPI_INITIALIZED (flag,ierr)
MPI_WTIME ()
MPI_WTICK ()
MPI_FINALIZE (ierr)
```

- Point-to-point communication functions
  - Message transfer from one process to another
- Collective communication functions
  - Message transfer involving all processes in a communicator



- MPI point-to-point operations typically involve message passing between two, and only two, different MPI tasks.
- One task is performing a send operation and the other task is performing a matching receive operation.
- There are different types of send and receive routines used for different purposes.
  - 1 Blocking send / blocking receive
  - 2 Non-blocking send / non-blocking receive
  - 3 Synchronous send

## Blocking vs. Non-blocking:

### ● Blocking send / receive

- send will "return" after it is safe to modify the application buffer (your send data) for reuse
- send can be synchronous i.e. handshake with the receive task to confirm a safe send.
- send can be asynchronous if a system buffer is used to hold the data for eventual delivery to the receive.
- receive only "returns" after the data has arrived and is ready for use by the program.

### ● Non-blocking send / receive

- behave similarly - they will return almost immediately.
- do not wait for any communication events to complete, such as message copying from user memory to system buffer space or the actual arrival of message.
- operations simply "request" the MPI library to perform the operation when it is able.

The user can not predict when that will happen.

- communications are primarily used to overlap computation with



# Blocking Message Passing Example

```
#include "mpi.h"
#include <stdio.h>

main(int argc, char *argv[]) {
 int numtasks, rank, dest, source, rc, count, tag=1;
 char inmsg, outmsg='x';
 MPI_Status Stat; // required variable for receive routines

 MPI_Init(&argc,&argv);
 MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
 MPI_Comm_rank(MPI_COMM_WORLD, &rank);

 // task 0 sends to task 1 and waits to receive a return message
 if (rank == 0) {
 dest = 1;
 source = 1;
 MPI_Send(&outmsg, 1, MPI_CHAR, dest, tag, MPI_COMM_WORLD);
 MPI_Recv(&inmsg, 1, MPI_CHAR, source, tag, MPI_COMM_WORLD, &Stat);
 }

 // task 1 waits for task 0 message then returns a message
 else if (rank == 1) {
 dest = 0;
 source = 0;
 MPI_Recv(&inmsg, 1, MPI_CHAR, source, tag, MPI_COMM_WORLD, &Stat);
 MPI_Send(&outmsg, 1, MPI_CHAR, dest, tag, MPI_COMM_WORLD);
 }

 // query receive Stat variable and print message details
 MPI_Get_count(&Stat, MPI_CHAR, &count);
 printf("Task %d: Received %d char(s) from task %d with tag %d \n",
 rank, count, Stat.MPI_SOURCE, Stat.MPI_TAG);

 MPI_Finalize();
}
```

```
program ping
include "mpi.h"

integer :: numtasks, rank, dest, source, count, tag, ierr
integer :: stat(MPI_STATUS_SIZE) ! required variable for receive routines
character :: inmsg, outmsg
outmsg = 'x'
tag = 1

call MPI_INIT(ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)
call MPI_COMM_SIZE(MPI_COMM_WORLD, numtasks, ierr)

! task 0 sends to task 1 and waits to receive a return message
if (rank .eq. 0) then
 dest = 1
 source = 1
 call MPI_SEND(outmsg, 1, MPI_CHARACTER, dest, tag, MPI_COMM_WORLD, ierr)
 call MPI_RECV(inmsg, 1, MPI_CHARACTER, source, tag, MPI_COMM_WORLD, stat,
 ierr)

 ! task 1 waits for task 0 message then returns a message
 else if (rank .eq. 1) then
 dest = 0
 source = 0
 call MPI_RECV(inmsg, 1, MPI_CHARACTER, source, tag, MPI_COMM_WORLD, stat, err)
 call MPI_SEND(outmsg, 1, MPI_CHARACTER, dest, tag, MPI_COMM_WORLD, err)
 endif

 ! query receive Stat variable and print message details
 call MPI_GET_COUNT(stat, MPI_CHARACTER, count, ierr)
 print *, 'Task ', rank, ': Received', count, 'char(s) from task', &
 stat(MPI_SOURCE), 'with tag', stat(MPI_TAG)

 call MPI_FINALIZE(ierr)

end program ping
```

# Blocking Message Passing Example

```
#include "mpi.h"
#include <stdio.h>

main(int argc, char *argv[]) {
 int numtasks, rank, right, left, buf[2], tag1=1, tag2=2;
 MPI_Request reqs[4]; // required variable for non-blocking calls
 MPI_Status stats[4]; // required variable for Waitall routine

 MPI_Init(&argc,&argv);
 MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
 MPI_Comm_rank(MPI_COMM_WORLD, &rank);

 // determine left and right neighbors
 left = rank-1;
 right = rank+1;
 if (rank == 0) left = numtasks-1;
 if (rank == (numtasks-1)) right = 0;

 // post non-blocking receives and sends for neighbors
 MPI_Irecv(&buf[0], 1, MPI_INT, left, tag1, MPI_COMM_WORLD, &reqs[0]);
 MPI_Irecv(&buf[1], 1, MPI_INT, right, tag2, MPI_COMM_WORLD, &reqs[1]);

 MPI_Isend(&rank, 1, MPI_INT, left, tag2, MPI_COMM_WORLD, &reqs[2]);
 MPI_Isend(&rank, 1, MPI_INT, right, tag1, MPI_COMM_WORLD, &reqs[3]);

 // wait for all non-blocking operations to complete
 MPI_Waitall(4, reqs, stats);

 printf("Task %d: Received from task %d with tag %d and from task %d with tag %d\n",
 rank, left, tag1, right, tag2);
 printf("Task %d: Send to task %d with tag %d and to task %d with tag %d\n",
 rank, left, tag2, right, tag1);

 MPI_Finalize();
}
```

```
program ringtopo
include 'mpif.h'

integer numtasks, rank, next, prev, buf(2), tag1, tag2, ierr, count
integer reqs(4) ! required variable for non-blocking calls
integer stats(MPI_STATUS_SIZE,4) ! required variable for WAITALL routine
tag1 = 1
tag2 = 2

call MPI_INIT(ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)
call MPI_COMM_SIZE(MPI_COMM_WORLD, numtasks, ierr)

! determine left and right neighbors
prev = rank-1
next = rank+1
if (rank .eq. 0) then
 prev = numtasks-1
endif
if (rank .eq. numtasks-1) then
 next = 0
endif

! post non-blocking receives and sends for neighbors
! Receive 1 from left and 2 from right
call MPI_Irecv(buf(1), 1, MPI_INTEGER, prev, tag1, MPI_COMM_WORLD, reqs(1), ierr)
call MPI_Irecv(buf(2), 1, MPI_INTEGER, next, tag2, MPI_COMM_WORLD, reqs(2), ierr)

! Send 1 to right and 2 to left
call MPI_Isend(rank, 1, MPI_INTEGER, prev, tag2, MPI_COMM_WORLD, reqs(3), ierr)
call MPI_Isend(rank, 1, MPI_INTEGER, next, tag1, MPI_COMM_WORLD, reqs(4), ierr)

! wait for all non-blocking operations to complete
call MPI_WAITALL(4, reqs, stats, ierr);

print '(S(a,i2))', 'Task ',rank,: 'Received from task', prev, ' with tag',tag1,
&
' and from task', next, ' with tag',tag2
print '(S(a,i2))', 'Task ',rank,: 'Send to task', prev, ' with tag',tag2, &
' and to task', next, ' with tag',tag1

! continue— do more work

call MPI_FINALIZE(ierr)

end program ringtopo
```

# Blocking Message Passing Example

```
[alp514.sol](1110): mpicc -o ringc ring.c
[alp514.sol](1113): srun -p eng -n 4 ./ringc
Task 0: Received from task 3 with tag 1 and from task 1 with tag 2
Task 0: Send to task 3 with tag 2 and to task 1 with tag 1
Task 1: Received from task 0 with tag 1 and from task 2 with tag 2
Task 1: Send to task 0 with tag 2 and to task 2 with tag 1
Task 2: Received from task 1 with tag 1 and from task 3 with tag 2
Task 2: Send to task 1 with tag 2 and to task 3 with tag 1
Task 3: Received from task 2 with tag 1 and from task 0 with tag 2
Task 3: Send to task 2 with tag 2 and to task 0 with tag 1
```

```
[alp514.sol](1111): mpif90 -o ringf ring.f90
[alp514.sol](1114): srun -p eng -n 4 ./ringf
Task 3: Received from task 2 with tag 1 and from task 0 with tag 2
Task 3: Send to task 2 with tag 2 and to task 0 with tag 1
Task 0: Received from task 3 with tag 1 and from task 1 with tag 2
Task 0: Send to task 3 with tag 2 and to task 1 with tag 1
Task 1: Received from task 0 with tag 1 and from task 2 with tag 2
Task 1: Send to task 0 with tag 2 and to task 2 with tag 1
Task 2: Received from task 1 with tag 1 and from task 3 with tag 2
Task 2: Send to task 1 with tag 2 and to task 3 with tag 1
```

## ● Tutorials

- ➊ MPI: <https://computing.llnl.gov/tutorials/mpi/>
- ➋ Advanced MPI:  
<https://hpc.llnl.gov/sites/default/files/DavidCronkSlides.pdf>
- ➌ CITutor: <https://www.citutor.org/>
- ➍ XSEDE HPC Monthly Workshop Series:  
<https://psc.edu/xsede-hpc-series-all-workshops>
- ➎ MPI Tutorial: <http://mpitutorial.com/>

## ● Books

- ➊ Beginning MPI (An Introduction in C) by Wesley Kendall
- ➋ Parallel Programming with MPI by Peter Pacheco (No relation)
- ➌ Using MPI - 2nd Edition: Portable Parallel Programming with the Message Passing Interface (Scientific and Engineering Computation) by William Gropp
- ➍ Parallel Programming in C with MPI and Openmp by Michael J. Quinn
- ➎ MPI: The Complete Reference by Marc Snir *et. al.*